

DIGITAL THREAT MONITORING PLATFORM

Master Strategy & MVP Launch Blueprint

Working concept: A security-as-a-service platform for SMBs that continuously discovers external digital exposure, detects suspicious threats, prioritizes risk, and explains what to fix.

Purpose

A practical document for the co-founder team to align on the problem, MVP scope, technology, validation, go-to-market, metrics, and 90-day execution plan.

Prepared: September 2026

1. Executive Summary

The opportunity is not to build another generic vulnerability scanner. The initial product should act as a **security autopilot for SMBs**: continuously monitor what a business exposes to the internet, identify meaningful security risks and suspicious brand activity, and turn technical findings into clear actions.

Core promise:

“Know what attackers can see about your business, understand what matters, and fix the highest-risk issues first.”

The MVP should deliberately stay narrow: **external asset discovery + exposure checks + threat/brand monitoring + risk prioritization + actionable alerts + AI-assisted explanations**. The objective is to prove recurring customer value before expanding into a full security platform.

2. Problem We Are Solving

- SMBs often have internet-facing assets they do not know about or actively monitor.
- Security tools can be technically powerful but difficult for non-security teams to interpret.
- One-time scans create reports, but customers need continuous awareness of what changed.
- Security teams face alert fatigue when tools produce too many low-confidence or low-value findings.
- Businesses need prioritized remediation, not a long list of technical findings.

Customer pain in one sentence: “I know cybersecurity is important, but I don't know what is exposed, what is actually dangerous, or what I should fix first.”

3. Ideal Customer Profile

Priority	Customer	Why they fit
1	SMB SaaS companies	Many internet-facing assets; strong need for uptime and trust; often limited security staff.
2	E-commerce businesses	Customer trust and brand impersonation matter; online assets are business-critical.
3	Digital/service companies	Need affordable security visibility without enterprise SOC tooling.
4	Small tech startups	Fast-changing infrastructure and limited dedicated security resources.

Initial ICP: companies with approximately 1–20 primary domains/assets, an online customer-facing presence, and no large internal security team.

4. Product Positioning

Do not position it as: “another vulnerability scanner.”

Position it as: “Continuous external security monitoring for businesses that don't have a dedicated security team.”

Long-term category: Digital Threat Monitoring / External Security Posture Management for SMBs.

5. MVP Scope

The MVP has five product pillars. Anything outside these should be treated as a later-stage feature unless customer interviews prove it is essential.

Module	MVP capability	Customer outcome
Asset Discovery	Verified domains, subdomains, DNS, certificates, IPs, technologies, public services.	Know the external attack surface.
Exposure Checks	Security headers, TLS/certificate checks, exposed services, technology/version signals, selected vulnerability intelligence.	Identify meaningful weaknesses.
Threat & Brand Monitoring	Suspicious domains, possible impersonation signals, relevant threat indicators and changes.	Detect threats targeting the brand.
Risk Engine	Severity + exposure + confidence + asset importance; finding lifecycle.	Know what to fix first.
AI Security Analyst	Explain findings, impact, evidence, recommended remediation and priority.	Make security understandable.

6. MVP User Journey

Step 1 — Sign up → create organization.

Step 2 — Verify ownership → DNS/file verification before active monitoring.

Step 3 — Add domain → system starts discovery.

Step 4 — Discover → build external asset inventory.

Step 5 — Analyze → run authorized checks asynchronously.

Step 6 — Prioritize → group duplicate findings and calculate risk.

Step 7 — Explain → AI converts evidence into plain-language guidance.

Step 8 — Monitor → scheduled jobs detect changes.

Step 9 — Alert → notify only meaningful events.

Step 10 — Resolve → customer fixes issue; next scan verifies status.

7. Product Architecture

Recommended architecture, optimized for a MERN/TypeScript development team:

Frontend: React/Next.js

API: Node.js + TypeScript

Database: MongoDB

Queue: Redis + BullMQ

Workers: isolated security discovery/scanning workers

AI: LLM layer constrained by verified finding evidence

Notifications: email first; Slack/WhatsApp later

Deployment: containerized services with separate API and worker scaling.

Critical design principle: scanning must be asynchronous. Never make the main API request wait for long-running discovery or analysis.

Flow: User → API → Job Queue → Security Worker → Evidence Store → Risk Engine → AI Explanation → Dashboard/Alert.

8. Core Data Model

Organization: owner, members, plan, settings

Asset: domain, subdomain, IP, service, technology, certificate, status

Scan: type, target, status, timestamps, worker metadata

Finding: asset, category, severity, confidence, evidence, impact, remediation, first_seen, last_seen, status

Threat: indicator, source, confidence, related asset/brand, timestamps

Alert: event, severity, delivery status, recipients

Audit Log: user/system action, timestamp, object, result.

9. Detection & Risk Strategy

The product's credibility will depend more on **accuracy and evidence** than on the number of checks.

Finding confidence: Confirmed / High / Medium / Low / Informational.

Never make a strong security claim from weak evidence. For example, detecting a software version does not automatically prove that a specific vulnerability is exploitable.

Recommended risk model: Risk = Severity × Exposure × Confidence × Asset Importance. Normalize the result to a customer-friendly 0–100 score.

Use a separate **business risk score** from raw vulnerability severity. This prevents a technically severe but non-exposed or low-confidence issue from dominating the dashboard.

10. False-Positive Control

- Store raw evidence for every finding.
- Use confidence scoring and evidence requirements.
- Deduplicate findings across repeated scans.
- Track first seen / last seen / resolved state.
- Allow customers to mark findings as accepted risk or false positive.
- Use human review during beta for high-impact detections.

North-star technical quality metric: percentage of high/critical findings that customers consider actionable and valid.

11. AI Strategy

AI should be a **reasoning and communication layer**, not the source of truth.

The model should receive structured evidence such as finding type, observed configuration, asset, confidence, references, and remediation guidance. It should not be allowed to invent evidence or claim exploitation without proof.

AI task	Allowed behavior
Explain	Translate technical findings into simple language.
Prioritize	Explain why an issue matters given verified context.
Remediate	Suggest practical defensive actions from approved guidance.
Summarize	Create executive security summaries and change reports.
Investigate	Ask clarifying questions or request more evidence instead of guessing.

12. Security & Abuse Controls

Because this is a cybersecurity product, authorization and safe operation are product requirements, not optional extras.

- Require customer asset/domain ownership verification.
- Only scan customer-authorized assets.
- Apply rate limits, scan budgets and worker isolation.
- Maintain audit logs for scans and sensitive actions.
- Build clear acceptable-use and abuse-reporting processes.
- Separate discovery/monitoring from any future active testing functionality.

- Protect stored credentials, API keys and customer data with encryption and least privilege.

13. 90-Day Execution Plan

Period	Focus	Deliverables	Success gate
Days 1–14	Validation + architecture	10–15 customer interviews, landing page, threat model, UX prototype, technical architecture.	At least 5 prospects confirm the problem is important and agree to beta.
Days 15–30	Foundation	Auth, organizations, domain verification, asset model, queue, dashboard shell.	A verified domain can be onboarded and jobs run reliably.
Days 31–50	Detection	Asset discovery, certificate/DNS/HTTP checks, technology signals, initial risk engine.	Repeatable scan with evidence and low obvious false positives.
Days 51–65	Threat monitoring	Suspicious-domain/brand monitoring, change detection, alerts.	Beta customers receive useful new-event alerts.
Days 66–80	AI + reporting	AI analyst, remediation summaries, security report, finding lifecycle.	Customers understand findings without security expertise.
Days 81–90	Beta → paid	10–20 active beta accounts, pricing, onboarding, support process, paid pilot.	At least 3–5 customers willing to pay or sign a paid pilot.

14. Team Responsibilities

Role	Primary ownership
Co-founder / Product	Customer interviews, ICP, roadmap, pricing, partnerships, onboarding and feedback loop.
Technical Lead	Architecture, backend, worker system, detection engine, security controls and reliability.
Frontend / Product Engineer	Dashboard, onboarding, findings UX, reports and organization management.
Security Research	Detection logic, threat intelligence sources, evidence quality and false-positive review.
Growth / Sales	Outbound, demos, beta recruitment, case studies and conversion.

For a small founding team, one person can own multiple roles. The key is that **product validation and technical execution happen in parallel**.

15. Go-to-Market Strategy

Phase 1 — Founder-led outbound

- Build a list of 100 target SMB SaaS/e-commerce companies.
- Offer a free external security snapshot or 14–30 day beta.
- Show 3–5 concrete findings with evidence rather than a generic sales pitch.
- Ask for permission to continuously monitor the verified domain.

Phase 2 — Product-led acquisition

- Free security exposure check as a lead magnet.
- Public educational content around phishing, exposed assets and security hygiene.
- Share anonymized security trend reports.

Phase 3 — Channel partnerships

- Digital agencies, MSPs and IT service providers can resell or manage security monitoring for their customers.

16. Pricing Hypothesis

Plan	Starting hypothesis	Purpose
Free	1 domain, limited monitoring	Acquisition and product discovery.
Starter	■999/month	Small businesses with basic monitoring needs.
Business	■2,999/month	Growing companies needing multiple assets and richer alerts.
Pro	■7,999+/month	Larger asset coverage, API/team/reporting features.

These prices are hypotheses, not final market pricing. Validate willingness-to-pay during beta. Price around **coverage and recurring value**, not simply number of scans.

17. Key Product Metrics

Metric	Why it matters
Verified domains onboarded	Measures activation.
Weekly active organizations	Measures recurring value.
New meaningful findings/week	Measures monitoring activity.
High/critical finding precision	Measures trust.
Alert open/action rate	Measures alert usefulness.
Time to first value	Measures onboarding quality.
Issues resolved after alert	Measures real customer impact.
30-day retention	Measures whether monitoring is valuable.
Free → paid conversion	Measures commercial viability.

North-star metric: number of verified customer organizations that actively monitor their assets and resolve at least one meaningful risk because of the platform.

18. What NOT to Build in MVP

- Full SIEM/SOC replacement.
- Complex enterprise compliance suite.
- Large mobile application.
- Dozens of integrations before product-market validation.
- Custom vulnerability database from scratch.
- Autonomous offensive exploitation.
- Huge threat-intelligence marketplace.
- Advanced endpoint detection and response.

The MVP should be narrow enough to launch in ~8–10 weeks and strong enough to prove recurring value.

19. Competitive Advantage / Moat

The scanning technology itself will not be the long-term moat. The defensible advantage should accumulate around:

- **Historical exposure graph:** know what changed over time.
- **High-quality evidence:** every finding is explainable and traceable.
- **Low false-positive rate:** fewer alerts, higher trust.
- **SMB-specific risk prioritization:** translate technical risk into business impact.
- **Remediation loop:** detect → alert → fix → verify → measure improvement.
- **Threat-to-asset correlation:** connect suspicious external signals to the customer's verified assets and brand.

The strongest long-term experience is not “Here are 500 vulnerabilities.” It is “**Here are the 3 things you should fix today, here is why, and here is proof when they are resolved.**”

20. Major Risks & Mitigation

Risk	Impact	Mitigation
False positives	Loss of trust	Evidence + confidence + human beta review.
Too broad MVP	Delayed launch	Strict five-module scope.
Unauthorized scanning	Legal/security risk	Ownership verification + authorization controls.
AI hallucination	Bad security advice	Evidence-grounded prompts and structured outputs.
High infrastructure cost	Poor margins	Async jobs, quotas, caching and staged scanning.
Weak customer willingness-to-pay	No business	Validate before building advanced features.
Alert fatigue	Users ignore product	Prioritize changes and high-confidence events.

21. Founder Operating Principles

- **Customer problem first:** every feature needs a customer reason.
- **Evidence over hype:** security claims must be defensible.
- **Accuracy over feature count:** 20 excellent detections beat 500 noisy ones.
- **Continuous value:** the product should get useful after onboarding, not become a static report.
- **Ship weekly:** release, measure, interview, improve.
- **Security by design:** customer data and authorization are core architecture concerns.

22. Immediate Next 7 Days

Day	Action	Output
1	Finalize ICP and problem statement	One-page product brief.
2	Interview 3–5 target customers	Pain points + current tools.
3	Interview 3–5 more + analyze competitors	MVP validation notes.
4	Create user flow and dashboard wireframe	Clickable/basic prototype.

Day	Action	Output
5	Finalize MVP detection list + data model	Technical specification.
6	Build domain verification + onboarding	First working vertical slice.
7	Set up queue/workers + first discovery job	End-to-end scan prototype.

23. Definition of MVP Success

The MVP is successful if, within the first 90 days:

- 10–20 businesses actively use it.
- Customers receive meaningful findings they did not already know about.
- False positives are low enough to maintain trust.
- Customers act on alerts and can verify remediation.
- At least 3–5 customers demonstrate willingness to pay.
- The founding team can explain the product's unique value in one sentence.

24. Final Product Thesis

We are not trying to replace a cybersecurity team. We are building the layer that gives a business a **continuous external security view** and tells it what deserves attention.

The winning loop is:

Discover → Verify → Analyze → Prioritize → Explain → Alert → Fix → Re-scan → Improve

If we can make that loop accurate, simple and recurring, we have the foundation for a much larger digital threat intelligence and security platform.

Co-founder decision checklist

- ■ Agree on ICP
- ■ Agree on MVP five modules
- ■ Assign technical/product ownership
- ■ Interview first 10 prospects
- ■ Finalize architecture
- ■ Start the first vertical slice
- ■ Define beta success metrics
- ■ Set a fixed 90-day launch target